

Twisted mass fermions on GPUs

An application to the LOEWE cluster at CSC

Lars Zeidlewicz

—

In collaboration with Matthias Bach, Owe Philipsen and
Christopher Pinke



Institut für Theoretische Physik
Goethe-Universität Frankfurt

Lattice QCD and high performance computing on GPUs
March 11, 2011

The plan of this talk


Introduction

- Finite temperature QCD
- Twisted mass lattice fermions

Status Report

- Solver
- Heatbath
- Using the LOEWE-CSC cluster


Physical objects of interest



GSI

properties of the
Quark-Gluon-Plasma

$$T > T_c$$





properties of the thermal
transition

$$T \approx T_c, \text{ possibly } \mu \neq 0$$

Chiral limit

What is the nature of the $N_f = 2$ chiral transition?




- Computational cost grows with inverse power of pion mass
- Infer chiral limit properties from universal scaling behaviour
- In any case: Small pion masses necessary

tmfT 2011

Lattice QCD at finite T

- Discretise spacetime by a hypercube of size $N_\sigma^3 \times N_\tau$ and lattice spacing a
- Temperature:

$$T = \frac{1}{aN_\tau}$$



- Carefully control cutoff effects:

$$F\left(\frac{1}{N_\tau}\right) = F_{\text{continuum}} + F_{(2)} \frac{1}{T^2 N_\tau^2} + \dots$$

- Need aspect ratio N_σ/N_τ as large as possible
 \Rightarrow large spatial volumes

Which lattice fermions?

Staggered fermions

- + Computationally cheap
- + Results at physical masses and below
- Theoretical doubts (“rooting”)

Wilson fermions

- + Theoretically sound
- + Improvement possible: clover & maximally twisted mass
- Explicitly broken chiral symmetry
 - ⇒ linear cutoff effects
 - ⇒ additive mass renormalisation


Chiral fermions

- + Theoretically sound
- + Lattice chiral symmetry
- Very expensive (non-local or 5 dimensional)

Twisted mass fermions: continuum

Frezzotti, Grassi, Sint, Weisz

Chiral symmetry transformation in $N_f = 2$ flavour space:



Creutz, PRD 76, 2007

$$\psi = \begin{pmatrix} u \\ d \end{pmatrix} \longrightarrow \chi = e^{-i\gamma_5 \tau^3 \omega/2} \psi$$


$$M \bar{\psi} \psi \longrightarrow \bar{\chi} \left(\underbrace{M \cos \omega}_{m_0} + i\gamma_5 \tau^3 \underbrace{M \sin \omega}_{\mu_0} \right) \chi$$

- Quark mass:

$$M_R = \sqrt{Z_m^2 m_0^2 + Z_\mu^2 \mu_0^2}$$

- Maximal twist:

$$m_0 = 0 \Rightarrow \omega = \pi/2$$




Twisted mass fermions: lattice

Frezzotti, Grassi, Sint, Weisz

fermion matrix with twisted mass term

$$M_{\text{tm}} = 1 + 2i\kappa a\mu_0\gamma_5\tau^3 - \kappa D_W[U]$$

- $N_f = 2$ mass-degenerate Wilson type fermions (inclusion of strange & charm possible)
- Hopping parameter:
 $\kappa = 1/(2am_0 + 8r)$
- Non-trivial bare parameter phase space (Aoki phase & Sharpe-Singleton plane)



tmfT, PRD 80, 2010

Fermion determinant

- Twisted mass acts as infrared regulator for the fermion matrix
- Flavour determinant can be evaluated explicitly:

$$\begin{aligned}\text{Det}_f M_{\text{tm}} &= \text{Det}_f \left(M_W[U] \mathbf{1}_f + 2i\kappa a\mu_0 \gamma_5 \tau^3 \right) \\ &= M_W[U] M_W^\dagger[U] + 4\kappa^2 (a\mu_0)^2\end{aligned}$$

- Original motivation for twisted mass term

Improvement

- Maximal twist is achieved for vanishing untwisted mass component
- Hopping parameter must be tuned to its critical value $\kappa_c(\beta)$
- Maximal twist ensures automatic $\mathcal{O}(a)$ improvement, i. e. absence of linear cutoff effects in physical observables

Frezzotti, Rossi, JHEP 0408

Example: Free dispersion relation

$$E(\mathbf{p}) = \sqrt{\mathbf{p}^2 + M^2} - a \frac{1}{2} \frac{M^3 \cos \omega}{\sqrt{\mathbf{p}^2 + M^2}} + \mathcal{O}(a^2)$$

tmLQCD package

- MPI-parallel CPU code for twisted mass fermions exists (freely available)
- C code developed within the European Twisted Mass Collaboration
Jansen, Urbach, *Comp. Phys. Commun.*, 180, 2009
- Algorithm: Mass preconditioning, multiple time-scale integrators, even-odd preconditioning. . .
Jansen et al. *Comp. Phys. Commun.*, 174, 2006
- Both $N_f = 2$ and $N_f = 2 + 1 + 1$ are implemented
- ETMC have used the two-flavour code down till 280 MeV pion masses
ETMC, *JHEP* 1008
- We test the code on the CPU part of LOEWE-CSC already

Motivation: LOEWE-CSC

- Top500 rank 22 in November 2010
(285 Tflops, rank 2 in Germany)
- Rank 8 of green IT machines
(741 Mflops/W, rank 4 in Germany)
- Approximately 800 nodes constitute a CPU/GPU hybrid structure
 - 2 AMD Opteron 12 core (\rightarrow 24 cores per node)
 - 1 AMD ATI Radeon 5870 (Cypress)
- Infiniband connected



CSC

Twisted mass simulations using OpenCL

Wishlist

- Have an OpenCL heatbath working
- Have a solver in OpenCL that can be applied to both CPU and GPU
- Have an OpenCL HMC
- Expand fermionic part to $N_f = 2 + 1 + 1$ PHMC
- Have a hybrid Hybrid Monte-Carlo, especially for the 24 core / 1 GPU architecture of a LOEWE-CSC node
- Ultimately: Have a program that can take advantage of the complete LOEWE-CSC cluster

Our Program

- New architecture & programming language
⇒ write code from scratch
- OpenCL is rapidly evolving
- Embed OpenCL code in C/C++ host environment

```
class opencl {  
public:  
  opencl(cl_device_type wanted){init(wanted);}  
  ~opencl(){finalize();}  
  ...  
}
```

Our Program

- New architecture & programming language
⇒ write code from scratch
- OpenCL is rapidly evolving
- Embed OpenCL code in C/C++ host environment

```
..  
#ifdef _RECONSTRUCT_TWELVE_  
collect_options<<' ' -D_RECONSTRUCT_TWELVE_'';  
#endif  
#ifdef _USEDDOUBLEPREC_  
collect_options<<' ' -D_USEDDOUBLEPREC_'';  
#endif  
...  
clerr = clBuildProgram(clprogram,1,  
                        &device,buildoptions.c_str(),0,0);
```

Our Program

- New architecture & programming language
⇒ write code from scratch
- OpenCL is rapidly evolving
- Embed OpenCL code in C/C++ host environment
- Compiler switches for single/double, 12-double reconstruction of $SU(3)$ matrices...
- Development started in December 2010
- OpenCL frontend class: ~ 2200 lines of C++ code
- Kernel code: $\mathcal{O}(3k)$ lines
(no BLAS libraries)

Solver in OpenCL

Status

- Target: Even-odd preconditioned solvers
BiCGStab, CG,...
- Achieved: Solver in OpenCL, no even-odd yet
- We have some experience from an existing CUDA solver
(F. Burger, HU Berlin)
- We have used the CUDA solver on the *gpu-scout* of CSC
(NVIDIA Teslas)


Current Limits

- For Cypress: RAM = 1 GB $\Rightarrow N_\sigma^3 \times N_\tau \leq 10^6$
e. g. $42^3 \times 12$, $38^3 \times 16$
- double precision, 12-double reconstruction, even-odd,
BiCGStab

Heatbath


Heatbath algorithm for $SU(3)$ pure gauge theory

- 12-double reconstruction for gauge fields (Clark et al.)
- For Cypress GPUs: 1 GB $\Rightarrow N_\sigma^3 \times N_\tau \leq 5 \cdot 10^6$
e.g. $76^3 \times 12$, $66^3 \times 16$
- Non-fermionic applications in our group:
QGP properties, hydrodynamics
- Use it independently and as first step towards an HMC



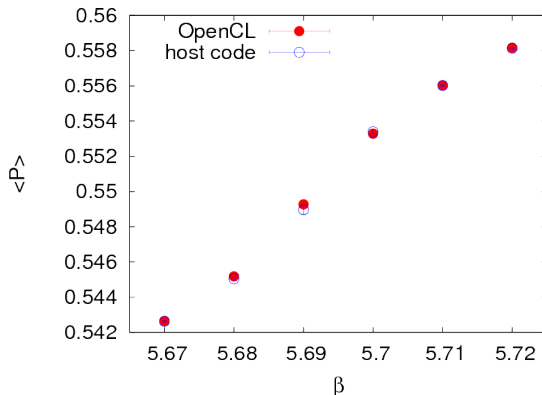
Checking the Heatbath

- Comparison to finite temperature results
Boyd et al. Nucl. Phys. B469, 1996
- $N_\tau = 8$, $N_\sigma = 16$, thermal transition at $\beta \approx 5.69$



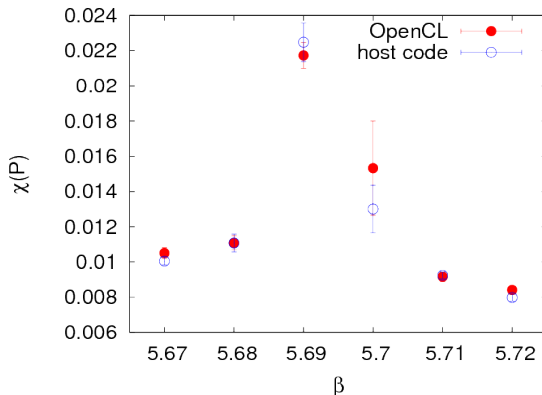
Checking the Heatbath

- Comparison to finite temperature results
Boyd et al. Nucl. Phys. B469, 1996
- $N_\tau = 8$, $N_\sigma = 16$, thermal transition at $\beta \approx 5.69$



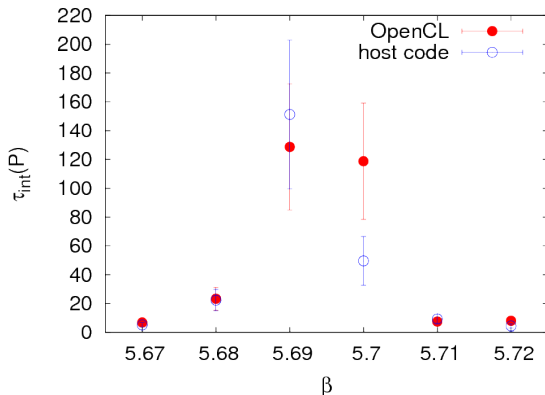
Checking the Heatbath

- Comparison to finite temperature results
Boyd et al. Nucl. Phys. B469, 1996
- $N_\tau = 8$, $N_\sigma = 16$, thermal transition at $\beta \approx 5.69$



Checking the Heatbath

- Comparison to finite temperature results
Boyd et al. Nucl. Phys. B469, 1996
- $N_\tau = 8$, $N_\sigma = 16$, thermal transition at $\beta \approx 5.69$



Hybrid OpenCL HMC

- One GPU is memory limited
- Try to distribute calculation on CPUs and attached GPU efficiently
LOEWE-CSC 24 cores & 1 GPU
- Distributing the calculation over several nodes leads to extra communication
- How can we make optimal use of the LOEWE-CSC cluster beyond one node?
- Possibly different algorithms are more suitable for that structure (e. g. DDHMC)

Conclusions and future plans

lots to do...

- Heatbath is ready, complete benchmarking
- Solver code needs further checks
- Probably, heavy performance optimisation possible
- Hybrid code must be implemented. We have not addressed questions concerning the global cluster structure yet

perspectives

- Fermion matrix can be exchanged
- OpenCL is programmed device independently

Thank you