

# **OpenCL based HMC Development Snapshot**

**Lars Zeidlewicz**

**Group Meeting – May 9, 2011**



**Institut für Theoretische Physik  
Goethe-Universität Frankfurt**

# ***Outline***

---

- **LOEWE-CSC**
- **Our code: desired features**
- **Code & project management**
- **Code overview**
- **Some words on OpenCL**

## Some numbers...

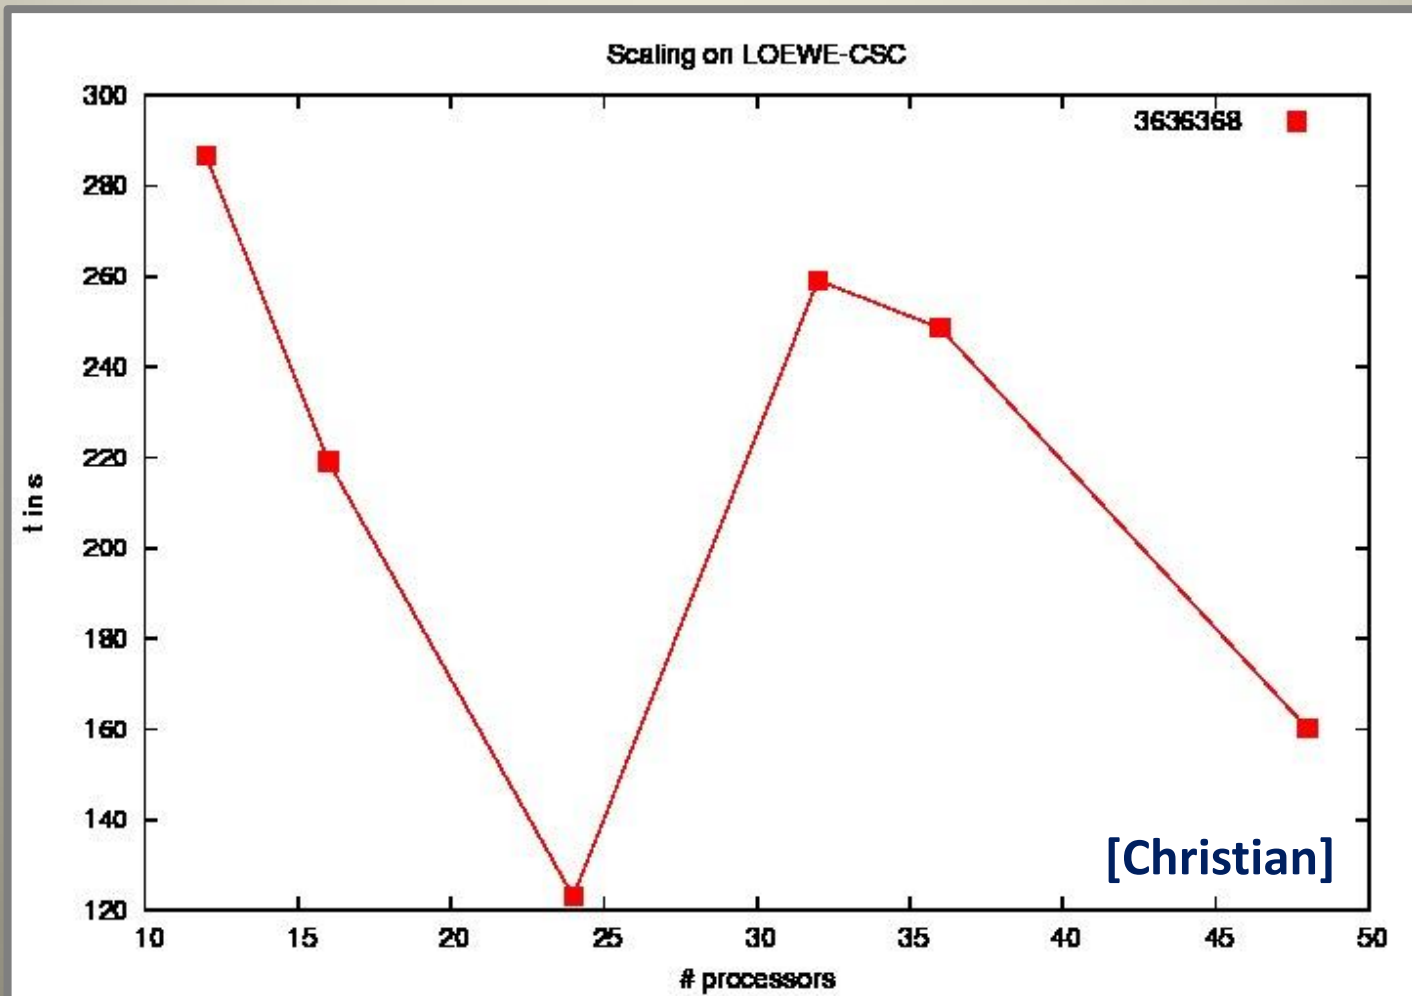
- **Top500 rank 22 (Nov 2010)**  
(285 Tflops, #2 in Germany)
- **Green500 rank 8**  
(741 Mflops/W, #4 in Germany)
- **About 800 nodes with**
  - 2 AMD Opteron 12 core
  - 1 AMD ATI Radeon 5870 (Cypress)

## **Status**

- **Testing phase**
- **Monthly CSC-admin meetings**
- **Current issues:**
  - Queue: Large jobs vs. single-core (Jobs wait for days and weeks)
  - Infiniband problems
  - 2 week downtime starts on Friday
- **GPUs so-far unused**

# LOEWE-CSC

## Status



## SLURM

- Admins hope to overcome queuing problems by switching to the SLURM scheduler

*(<https://computing.llnl.gov/linux/slurm/slurm.html>)*

- Today: 29 nodes  
After downtime: 100 nodes
- OpenMPI tested (*only!*)

## SLURM

```
#!/bin/sh
#
#SBATCH --nodes=4
#SBATCH --time=1
#SBATCH --cpus-per-task=1
#SBATCH --job-name=testjob
#
OMP_NUM_THREADS=1
Srun -ln4 --mpi=none hostname
```

- How does that work for a "real" MPI program?
- SLURM can be loaded as a module:  
    module load slurm  
    module show slurm  
    (sinfo, squeue,...)
- Working examples have been promised

- OpenMPI tested (*only!*)

# Code Features

---

## Physics

- SU(3) heatbath
- Fermion inverter  
(different actions)
- Hybrid Monte-Carlo
- *Chemical potential*
- *Smearing*
- ...



# *Code Features*

---

## Programming

- Hybrid structure
  - *Within nodes*
  - *Beyond one node*
- OOP for host code
- Large flexibility
  - Switches (single/double etc)
  - User defined types, access functions to facilitate changes

# Code Features

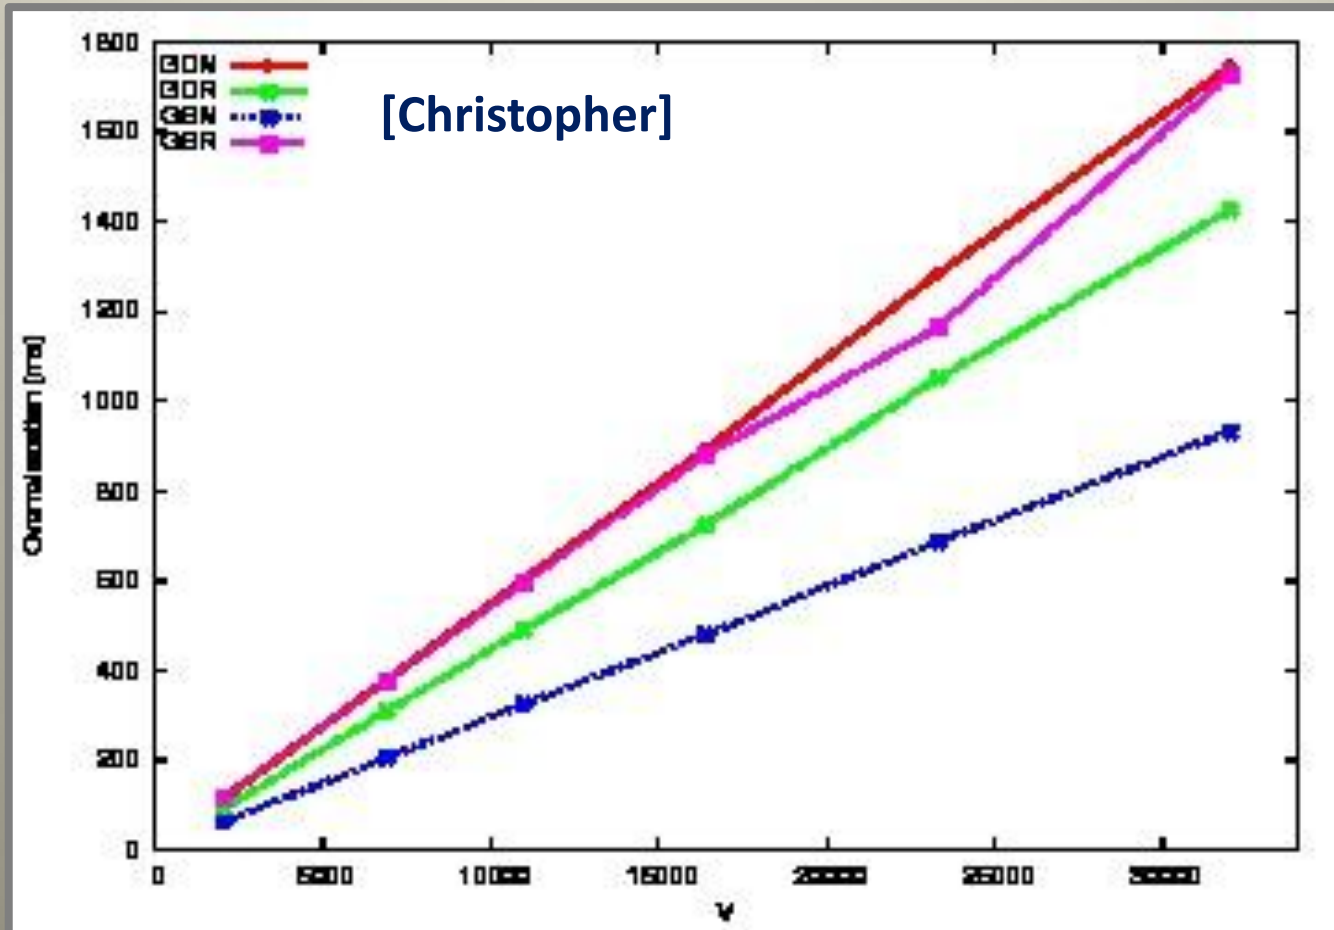
---

## GPUs

- **Today:**
  - Memory Management [Matthias]
  - Benchmarking
  - Workgroup sizes [Matthias]
- **Future:**
  - Several devices

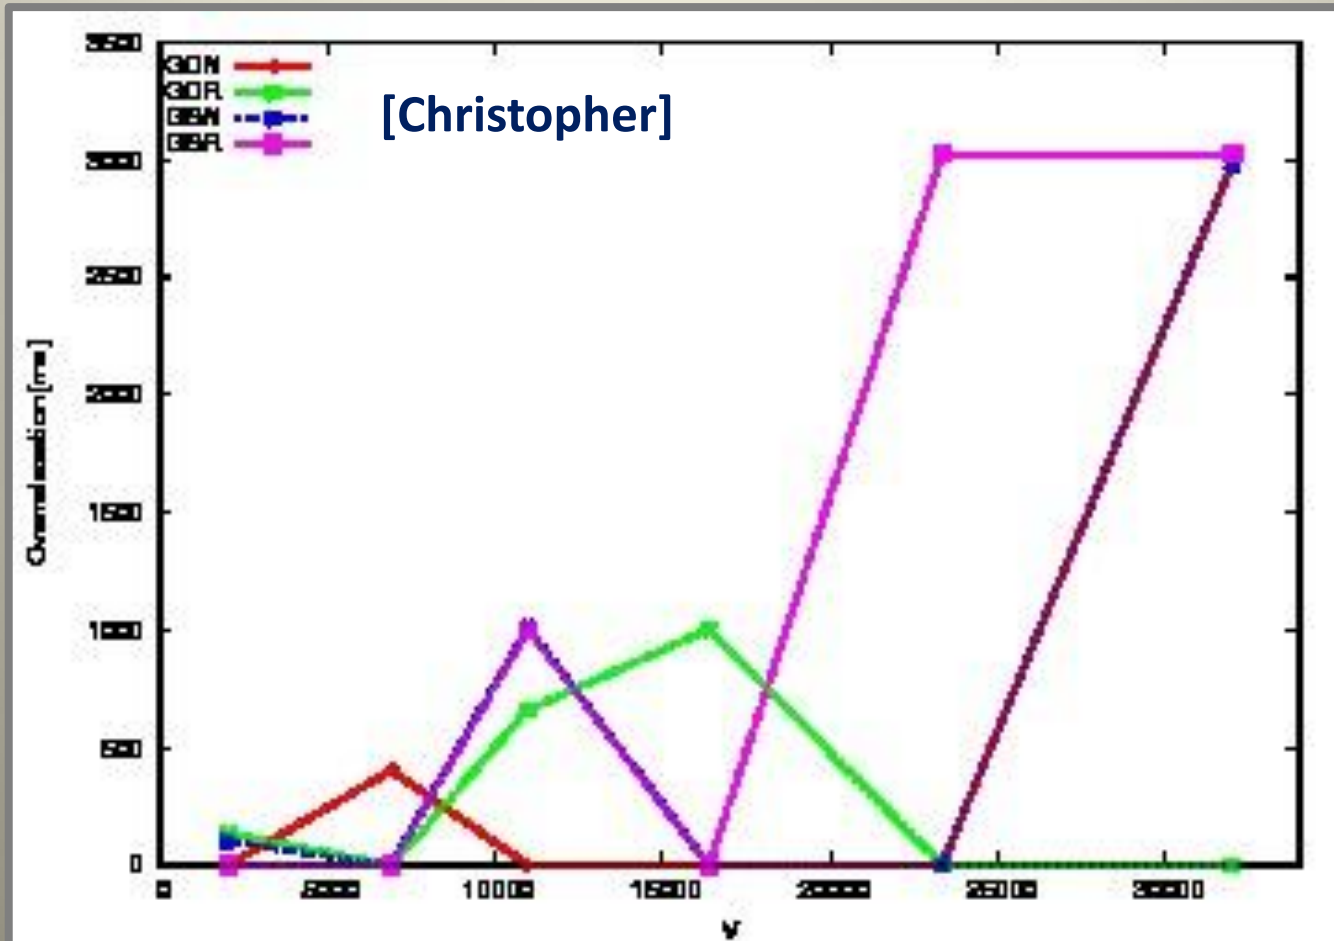
# Code Features

## GPUs benchmark – gpu-dev



# Code Features

## GPUs benchmark – loewe



# Code & Project Management

## git

- Code is subject to version management by git

```
lars@rupert: ~/clhmc_dev/clhmc
Datei Bearbeiten Ansicht Terminal Hilfe
commit 51b5ff5aa8f4402db0aa4691acf8f1e7e30aa615
Author: Lars Zeidlewicz <zeidlewicz@th.physik.uni-frankfurt.de>
Date: Wed May 4 13:07:42 2011 +0200

    fixed some inconsistencies for heatbath, correct names for saved configs from gaugefield class

commit a15e317d3edc05449cc1f637caa98e9dbd7e86aa
Author: Matthias Bach <bach@compeng.uni-frankfurt.de>
Date: Tue May 3 18:31:33 2011 +0200

    Initialize device random number generator from a large seed stored in a file

commit 50bd8d076d37183798d79d46286fca922c5e99b2
Author: Matthias Bach <bach@compeng.uni-frankfurt.de>
Date: Tue May 3 18:00:17 2011 +0200

    Source format

commit cd2e862174bb031f3efe597e0e9b291be5f739ef
Author: Lars Zeidlewicz <zeidlewicz@th.physik.uni-frankfurt.de>
Date: Tue May 3 17:58:05 2011 +0200

    missing gaugefield.cpp from my previous commit

commit 2f823e5c1239ef33d7268755024cd17ec0ed2f40
Author: Matthias Bach <bach@compeng.uni-frankfurt.de>
Date: Tue May 3 17:43:36 2011 +0200

    Fixed broken merge of opencl.h

commit 7ae28590a4f702902a924d022031f9373055ff92
Author: Christian Schäfer <cschaefer@th.physik.uni-frankfurt.de>
```

# Code & Project Management

## Online repository

[Matthias]

- <http://code.compeng.uni-frankfurt.de/projects/clhmc>

The screenshot shows a web browser displaying the project page for 'OpenCL based HMC' on the 'code.compeng.uni-frankfurt.de' website. The page has a blue header with navigation links: 'Hauptseite', 'Meine Seite', 'Projekte', and 'Hilfe'. A search bar is present on the right side of the header. Below the header, there are tabs for 'Übersicht', 'Aktivität', 'Tickets', 'Neues Ticket', 'Projektarchiv', and 'Konfiguration'. The 'Übersicht' tab is selected, showing a summary of the project. The main content area is divided into three columns. The left column contains a 'Tickets' section with a list of open tickets: 'Bug: 8 offen / 8', 'Feature: 7 offen / 9', 'Unit Test: 0 offen / 0', 'Benchmark: 0 offen / 0', and 'Test Setup: 0 offen / 0'. Below this is a 'Git Repository' section with 'Read Only URL' and 'Developer URL' for the project. The middle column contains a 'Mitglieder' (Members) section listing the project manager (Matthias Bach) and developers (Christian Schäfer, Christopher Pinke, Lars Zeidlewicz, Matthias Bach, Michael Fromm, Stefano Lottini). The right column contains an 'Aufgewendete Zeit' (Time Spent) section showing '18.25 Stunden' and a link to 'Details | Bericht'.

OpenCL based HMC

Suche:  OpenCL based HMC

Übersicht Aktivität Tickets Neues Ticket Projektarchiv Konfiguration

**Übersicht**

We want to write a complete HMC simulation programme with Wilson-type quarks (twisted mass action) using OpenCL for a hybrid structure consisting of CPUs and GPUs.

**Tickets**

- Bug: 8 offen / 8
- Feature: 7 offen / 9
- Unit Test: 0 offen / 0
- Benchmark: 0 offen / 0
- Test Setup: 0 offen / 0

Alle Tickets anzeigen | Kalender | Gantt-Diagramm

**Git Repository**

Read Only URL:

- <git://code.compeng.uni-frankfurt.de/clhmc.git>

Developer URL:

- <git@code.compeng.uni-frankfurt.de:clhmc.git>

**Mitglieder**

Manager: Matthias Bach

Developer: Christian Schäfer, Christopher Pinke, Lars Zeidlewicz, Matthias Bach, Michael Fromm, Stefano Lottini

Reporter: Christian Schäfer, Christopher Pinke, Lars Zeidlewicz, Matthias Bach, Michael Fromm, Stefano Lottini

**Aufgewendete Zeit**

18.25 Stunden

Details | Bericht

# Code & Project Management

## Online repository

[Matthias]

- <http://code.compeng.uni-frankfurt.de/projects/clhmc>

The screenshot shows a web browser displaying the project page for 'OpenCL based HMC' on the website 'code.compeng.uni-frankfurt.de/projects/clhmc'. The page has a blue header with navigation links: 'Hauptseite', 'Meine Seite', 'Projekte', and 'Hilfe'. Below the header, there's a search bar and a dropdown menu set to 'OpenCL based HMC'. The main content area is divided into sections: 'Übersicht' (Overview), 'Aktivität', 'Tickets', 'Neues Ticket', 'Projektarchiv', and 'Konfiguration'. The 'Übersicht' section contains a description of the project: 'We want to write a complete HMC simulation program for heavy quarks (twisted mass action) using OpenCL for a hybrid structure consisting...'. The 'Tickets' section lists several tickets: 'Bug: 8 offen / 8', 'Feature: 7 offen / 9', 'Unit Test: 0 offen / 0', 'Benchmark: 0 offen / 0', and 'Test Setup: 0 offen / 0'. The 'Mitglieder' (Members) section lists the Manager 'Matthias Bach' and Developers 'Christian Schäfer', 'Christopher Pinke', 'Lars Zeidlewicz', 'Matthias Bach', 'Michael', and 'Sebastian Löffel'. The 'Aufgewendete Zeit' (Time spent) section shows '18.25 Stunden' and links to 'Details' and 'Bericht'. The 'Git Repository' section provides 'Read Only URL' and 'Developer URL' for the project. An orange callout box points to the 'Tickets' section with the text 'Side remark: The name?!?'.

Side remark: The name?!?

- OpenCL based HMC
- OpTiMaL

# ***Code & Project Management***

---

## **Common style**

[Matthias]

- **Matthias has suggested to enforce a common coding style via**

```
astyle --options=astylerc <FILENAME>
```

- **To be discussed...**



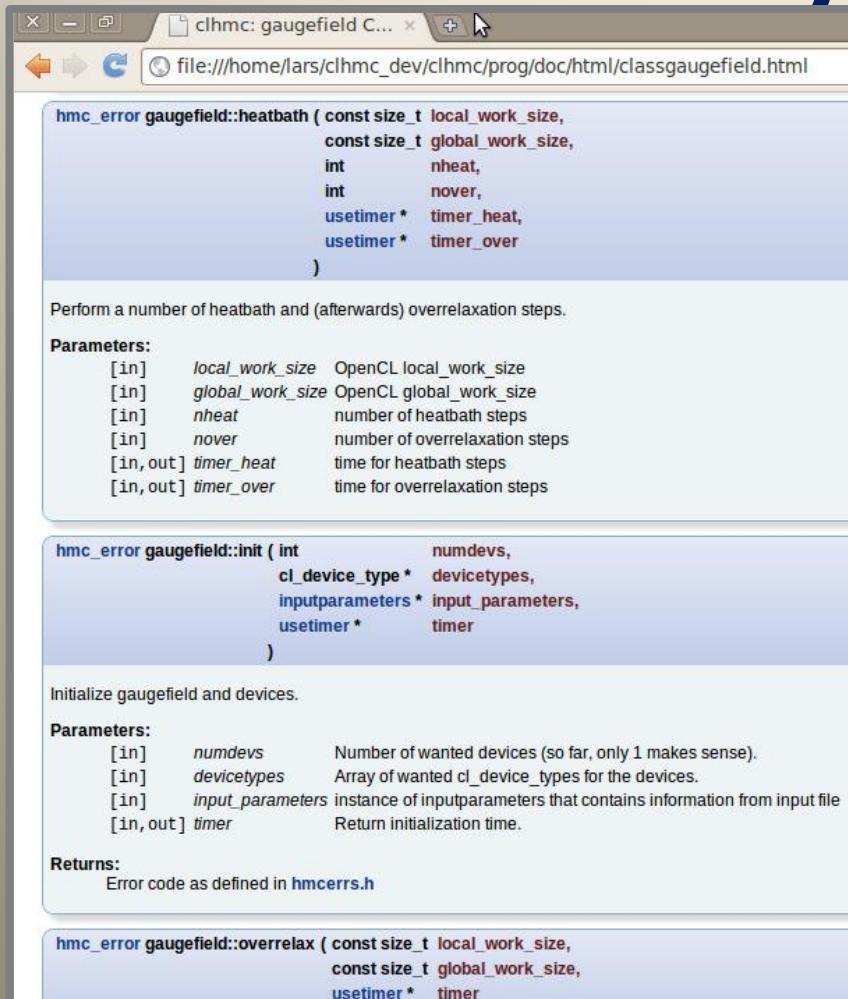
# Code & Project Management

## Doxygen

```
File Edit Options Buffers Tools C Help
class gaugefield {
public:
/**errs.h
 * Initialize gaugefield and devices.
 * @param[in] numdevs Number of wanted devices (so far, only 1 makes sense).
 * @param[in] devicetypes Array of wanted cl_device_types for the devices.
 * @param[in] input_parameters instance of inputparameters that contains information from input file
 * @param[in,out] timer Return initialization time.
 * @return Error code as defined in hmcerrs.h
 */
hmc_error init(int numdevs, cl_device_type* devicetypes, inputparameters* input_parameters, usetimer* timer);
/**keCache.txt cmake install.cmake Doxyfile
 * Free gaugefield and device allocations.
 */
hmc_error finalize();
/**keCache.txt cmake install.cmake Doxyfile
 * Save gaugefield to file.
 */
hmc_error save(int number);
/**e[3]: execvp: ./Doxyfile: Keine Berechtigung
 * Copy gaugefield to devices (currently: to device).
 * @param[in,out] timer copy-time
 */
hmc_error copy_gaugefield_to_devices(usetimer* timer);
/**keCache.txt cmake install.cmake Doxyfile
 * Copy random array to devices (currently: to device).
 * @param[in,out] timer copy-time
 */
hmc_error copy_rndarray_to_devices(hmc_rndarray host_rndarray, usetimer* timer);
/**
 * Copy random array from devices (currently: from device).
 * @param[in,out] timer copy-time
 */
hmc_error copy_rndarray_from_devices(hmc_rndarray rndarray, usetimer* timer);
/**
-UU-:----F1 gaugefield.h 14% L62 Git-master (C/l Abbrev)-----
```

# Code & Project Management

## Doxygen



The screenshot shows a web browser window with the address bar displaying 'file:///home/lars/clhmc\_dev/clhmc/prog/doc/html/classgaugefield.html'. The main content area shows the Doxygen documentation for the 'heatbath' function. The function signature is: `hmc_error gaugefield::heatbath ( const size_t local_work_size, const size_t global_work_size, int nheat, int nover, usetimer * timer_heat, usetimer * timer_over )`. Below the signature, there is a description: 'Perform a number of heatbath and (afterwards) overrelaxation steps.' followed by a 'Parameters:' section with a table of arguments. The table lists: `local_work_size` (OpenCL local\_work\_size), `global_work_size` (OpenCL global\_work\_size), `nheat` (number of heatbath steps), `nover` (number of overrelaxation steps), `timer_heat` (time for heatbath steps), and `timer_over` (time for overrelaxation steps). Below this, the 'init' function is shown with signature: `hmc_error gaugefield::init ( int numdevs, cl_device_type * devicetypes, inputparameters * input_parameters, usetimer * timer )`. Its description is 'Initialize gaugefield and devices.' followed by a 'Parameters:' section with a table: `numdevs` (Number of wanted devices), `devicetypes` (Array of wanted cl\_device\_types), `input_parameters` (instance of inputparameters), and `timer` (Return initialization time). A 'Returns:' section indicates 'Error code as defined in hmcerrs.h'. At the bottom, the 'overrelax' function signature is partially visible: `hmc_error gaugefield::overrelax ( const size_t local_work_size, const size_t global_work_size, usetimer * timer`.

```
hmc_error gaugefield::heatbath ( const size_t local_work_size,
                                const size_t global_work_size,
                                int nheat,
                                int nover,
                                usetimer * timer_heat,
                                usetimer * timer_over
                                )

Perform a number of heatbath and (afterwards) overrelaxation steps.

Parameters:
[in]    local_work_size  OpenCL local_work_size
[in]    global_work_size OpenCL global_work_size
[in]    nheat            number of heatbath steps
[in]    nover            number of overrelaxation steps
[in, out] timer_heat     time for heatbath steps
[in, out] timer_over     time for overrelaxation steps

hmc_error gaugefield::init ( int numdevs,
                             cl_device_type * devicetypes,
                             inputparameters * input_parameters,
                             usetimer * timer
                             )

Initialize gaugefield and devices.

Parameters:
[in]    numdevs          Number of wanted devices (so far, only 1 makes sense).
[in]    devicetypes      Array of wanted cl_device_types for the devices.
[in]    input_parameters instance of inputparameters that contains information from input file
[in, out] timer          Return initialization time.

Returns:
Error code as defined in hmcerrs.h

hmc_error gaugefield::overrelax ( const size_t local_work_size,
                                  const size_t global_work_size,
                                  usetimer * timer
```

Use doxygen  
format for  
comments

# ***Code & Project Management***

---

## **cmake**

- **Nice build-tool**
- **Have a look at  
clhmc/prog/INSTALL**
- **Creates Makefile:**
  - make hmc
  - make heatbath

## **Overview**

- **Three different types of code:**
  - **Host management code**  
(→ C++)
  - **Host functional code**  
(deprecated?)
  - **Device functional code**  
(→ OpenCL)

## Heatbath

- **Based on Christopher's code**
- **Check: Boyd et al.:**

Nucl. Phys. B469 (1996) 419-444

- **Host and device versions**

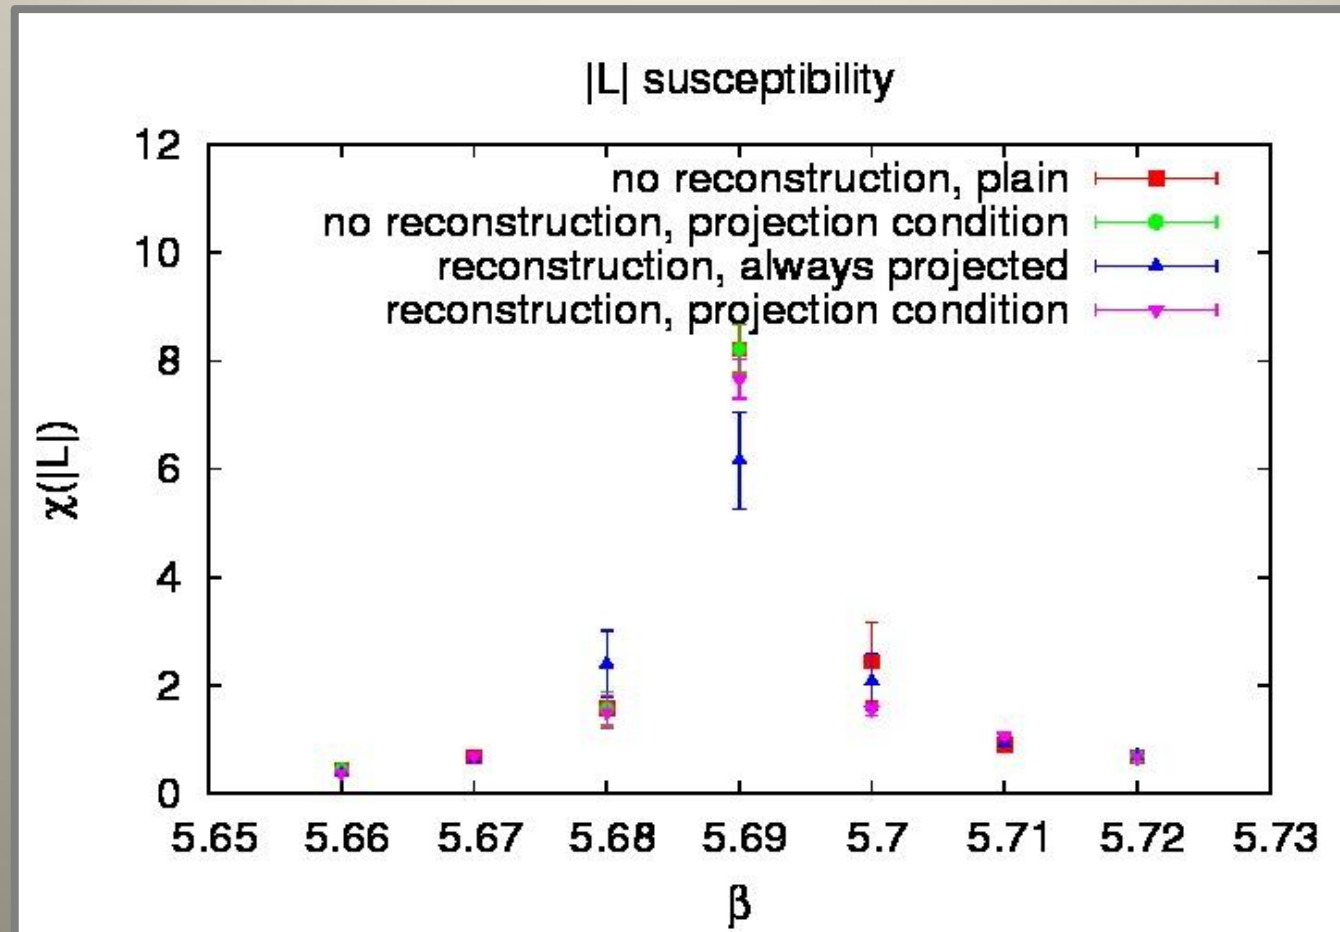
still to check device version with final status

- **Current improvement:**

Implement an optimal projection to  $SU(3)$   
[Christian]

# Code status

## Heatbath



# *Code status*

---

## HMC & inverter

- **Work in progress** [Christopher, Steo]
- **Host code inverter needs to be checked** [Lars]
- **So-far host code only**
- **Molecular dynamics force?** [Christopher]

## Management

- `int main() {}` in
  - `hmc.cpp`  
(mixes all progs by compiler flags)
  - `heatbath.cpp`
  - ...
- Management in classes:
  - `gaugefield`
  - `opencl`
  - ...



# *Code status*

---

## globaldefs.h

- NC, NSPIN, NDIM,...
- Right now: NT, NS
- Right now: work\_sizes
- Derived constants (VOLSPACE...)
- Available in all of the code  
(sometimes useful:  
kernel flag `_INKERNEL_`)

# *Code status*

---

## types.h

- `hmc_float`
- `hmc_complex`
- `hmc_su3matrix`
- ...
- Switches: `INKERNEL`,  
`USEDDOUBLEPREC`,  
`RECONSTRUCT_TWELVE`

## inputparameters

- Class `inputparameters`
- Reads input file
- Look at method  
`set_defaults()`

```
kappa=.125  
mu=0.06  
cgmax=1000  
beta=6  
#cgmax=3  
prec=64  
startcondition=cold  
sourcefile=conf.save  
thermalizationsteps=5  
heatbathsteps=1000
```

# *Code status*

---

## gaugefield.h

- ILDG storage format
- Gauge field is really the central object:
  - hmc trajectories
  - heatbath
  - inversion
- Gaugefield class: the place to manage inter-device communication

# Code status

## openc1.h

- Wraps all device operations
- `class openc1`

openc1 Class Reference	
#include <openc1.h>	
List of all members.	
Public Member Functions	
	<code>openc1 (cl_device_type wanted, const size_t ls, const size_t gs, usetimer *timer, inputparameters *parameters)</code>
	<code>openc1 ()</code>
<code>hmc_error</code>	<code>init (cl_device_type wanted_device_type, const size_t local_work_size, const size_t global_work_size, usetimer *timer, inputparameters *parameters)</code>
<code>hmc_error</code>	<code>copy_gaugefield_to_device (hmc_gaugefield *host_gaugefield, usetimer *timer)</code>
<code>hmc_error</code>	<code>copy_rndarray_to_device (hmc_rndarray host_rndarray, usetimer *timer)</code>
<code>hmc_error</code>	<code>copy_rndarray_from_device (hmc_rndarray rndarray, usetimer *timer)</code>
<code>hmc_error</code>	<code>get_gaugefield_from_device (hmc_gaugefield *host_gaugefield, usetimer *timer)</code>
<code>hmc_error</code>	<code>run_heatbath (hmc_float beta, const size_t local_work_size, const size_t global_work_size, usetimer *timer)</code>
<code>hmc_error</code>	<code>run_overrelax (hmc_float beta, const size_t local_work_size, const size_t global_work_size, usetimer *timer)</code>
<code>hmc_error</code>	<code>gaugeobservables (const size_t local_work_size, const size_t global_work_size, hmc_float *plaq, hmc_float *tplaq, hmc_float *splaq, hmc_complex *pol, usetimer *timer1, usetimer *timer2)</code>
<code>hmc_error</code>	<code>finalize ()</code>
Public Attributes	
	<code>std::vector&lt; std::string &gt; cl_kernels_file</code>

# *Code status*

---

## Random numbers

- Random numbers on GPU are tricky; avoid cross correlations between threads
- Vectorized, parallel RNG?!?
- Right now: NR generator
- Implement RanLux [Matthias]

## Operations

- Implement operations for different types:

`host_operations_complex`  
`host_operations_fermionmatrix`  
`host_operations_gaugefield`  
`host_operations_spinor`  
`host_operations_spinorfield`

- Maybe we can get that a bit more sorted?

[currently: Steo, Christian?]

# Code status

## Testing function

- Deprecated?

```
void testing_correlator(hmc_gaugefield* gf, inputparameter  
void testing_fermionmatrix();  
void testing_eoprec_spinor();  
void print_fullspinorfield(hmc_spinor* in);  
void testing_spinor();  
void print_su3mat(hmc_su3matrix* A);  
void print_staplemap(hmc_staplemap* A);  
void print_linkplusstaplemap(hmc_gaugefield * in, int  
void testing_su3mat();  
void testing_gaugefield();  
void testing_geometry();  
void testing_su3matrix(hmc_gaugefield * in, int spacepos,  
void testing_adjoin(hmc_gaugefield * in, int spacepos, in  
void testing_det_global(hmc_gaugefield * in);  
void testing_matrix_ops(hmc_gaugefield * in);  
  
void testing_heatbath_norandommat_no123(hmc_su3matrix * s  
void testing_heatbath_no123(hmc_su3matrix * su3_in, hmc_s  
y, int * cter_out);  
  
void testing_heatbath(hmc_su3matrix * su3_in, hmc_staple  
* cter_out);  
  
void testing_colorvector_ops();  
void testing_matrix_spinor_ops();  
void testing_matrix_spinor_functions();  
void testing_fermionmatrix_functions();  
  
void unit_spinor(hmc_spinor * in);  
void i_spinor(hmc_spinor * in);  
void print_spinor(hmc_spinor * in);  
void set_comp_to_one_spinor(hmc_spinor * in, int comp);
```



# *Code status*

---

## Compile time flags

- `USEDDOUBLEPREC ,  
USE_EOPREC`
- `RECONSTRUCT_TWELVE`
- `FERMIONS , HMC ,  
BENCHMARKS`

➔ Untangle into different programs

## SU(3) reconstruct

- M. Clark et al.

Comput. Phys. Commun., 181:1517-1528, 2010

$$\begin{pmatrix} \mathbf{a} \\ \mathbf{b} \\ \mathbf{c} \end{pmatrix} = \begin{pmatrix} a_1 & a_2 & a_3 \\ b_1 & b_2 & b_3 \\ c_1 & c_2 & c_3 \end{pmatrix}$$

$$\mathbf{c} = (\mathbf{a} \times \mathbf{b})^*$$

- `hmc_complex`  
`reconstruct_su3(hmc_su3matrix *in,`  
`int j);`

## Device code

- `.cl` files read in by `opencl.cpp`
- Compilation at runtime
- `cl_kernelsfiles.cl` for debug info

## Device code

```
/** @file
 * Inclusion and definition of types and definitions required in the Device code.
 */
//opengl_header.cl

#ifdef USEDOUBLEPREC
#pragma OPENCL EXTENSION cl_amd_fp64 : enable
//#pragma OPENCL EXTENSION cl_khr_fp64 : enable
#endif

#include "globaldefs.h" //NDIM, NSPIN, NC
#include "types.h"

//!!!CP: why is this here?
// #define VOLSPACE NSPACE*NSPACE*NSPACE
#define VOL4D VOLSPACE*NTIME

//for hmc ocl su3matrix
#ifdef RECONSTRUCT_TWELVE_
#define SU3SIZE NC*(NC-1)
#define STAPLEMATRIXSIZE NC*NC
#else
#define SU3SIZE NC*NC
```

- [.cl files read in by opengl.cpp](#)
- [compilation at runtime](#)  
(new possibilities)
- [look also at cl\\_kernelsfiles.cl](#)

## Some remarks...

- Build-log
- Kernels, kernel args
- Pragmas (single, double)
- Optimise worksizes [Matthias...]

## Some remarks...

```
init device #0
OpenCL being initialized...
  CL_PLATFORM_NAME:      ATI Stream
  CL_PLATFORM_VENDOR:    Advanced Micro Devices, Inc.
  CL_PLATFORM_VERSION:   OpenCL 1.1 ATI-Stream-v2.3 (451)

  1 device of wanted type has been found.
  Device information:
    CL_DEVICE_NAME:      Intel(R) Core(TM) i5 CPU          M 430  @ 2.27GHz
    CL_DEVICE_VENDOR:    GenuineIntel
    CL_DEVICE_TYPE:      CPU
    CL_DEVICE_VERSION:   OpenCL 1.1 ATI-Stream-v2.3 (451)
    CL_DEVICE_EXTENSIONS: cl_amd_fp64 cl_khr_global_int32_base_atomics cl_khr_global_int32_extended_atomics
cl_khr_local_int32_base_atomics cl_khr_local_int32_extended_atomics cl_khr_int64_base_atomics cl_khr_int64_extended_atom
ics cl_khr_byte_addressable_store cl_khr_gl_sharing cl_ext_device_fission cl_amd_device_attribute_query cl_amd_media_ops
  cl_amd_popcnt cl_amd_printf
Create context...
Create command queue...
Read kernel source from file: /home/lars/clhmc_dev/clhmc/prog/opencl_header.cl
Read kernel source from file: /home/lars/clhmc_dev/clhmc/prog/opencl_geometry.cl
Read kernel source from file: /home/lars/clhmc_dev/clhmc/prog/opencl_random.cl
Read kernel source from file: /home/lars/clhmc_dev/clhmc/prog/opencl_operations_complex.cl
Read kernel source from file: /home/lars/clhmc_dev/clhmc/prog/opencl_operations_gaugefield.cl
Read kernel source from file: /home/lars/clhmc_dev/clhmc/prog/opencl_update_heatbath.cl
Read kernel source from file: /home/lars/clhmc_dev/clhmc/prog/opencl_gaugeobservables.cl
Create program...
Build program...
  build options: -D_INKERNEL_ -DNSPACE=8 -DNTIME=4 -DVOLSPACE=512 -DSPINORSIZE=12 -DHALFSPINORSIZE=6 -DSPINORFIEL
DSIZE=24576 -DEOPREC_SPINORFIELDSIZE=12288 -D_USEDDOUBLEPREC_ -I/home/lars/clhmc_dev/clhmc/prog
finished building program
Build Log:
Create buffer for gaugefield...
Create buffer for random numbers...
Create buffer for gaugeobservables...
Create heatbath kernels...
Create gaugeobservables kernels...
writing gaugefield to lime file
```

# *That's it*

---

## To do

- Gaugefield class [Lars]
- Inter-device strategy
- Device code (esp. hmc) [Christopher]
- Memory management [Matthias]
- OpenSource, GPL?